



AVL Trees

Invented in 1962 by
G. M. Adelson-Velskii
E. M. Landis




AVL Tree Definition

- ◆ AVL trees are balanced
- ◆ An AVL Tree is a **binary search tree** such that for every internal node v of T , the **heights of the children of v can differ by at most 1**



An example of an AVL tree where the heights are shown next to the nodes:

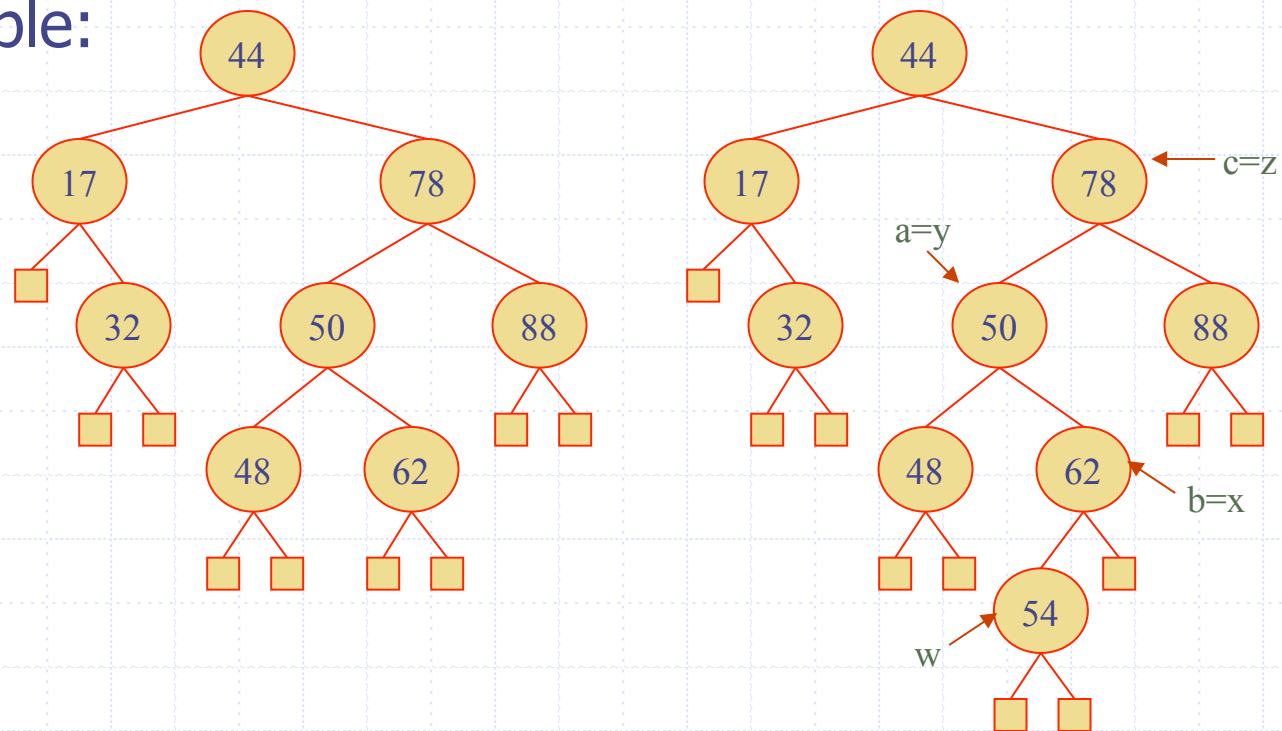
Height of an AVL Tree



- ◆ **Fact:** The height of an AVL tree storing n keys is $O(\log n)$.
- ◆ **Proof:** Let us bound $n(h)$: the minimum number of internal nodes of an AVL tree of height h .
- ◆ We easily see that $n(1) = 1$ and $n(2) = 2$
- ◆ For $n > 2$, an AVL tree of height h contains the root node, one AVL subtree of height $n-1$ and another of height $n-2$.
- ◆ That is, $n(h) = 1 + n(h-1) + n(h-2)$
- ◆ Knowing $n(h-1) > n(h-2)$, we get $n(h) > 2n(h-2)$. So
 - $n(h) > 2n(h-2)$, $n(h) > 4n(h-4)$, $n(h) > 8n(h-6)$, ... (by induction),
 - $n(h) > 2^i n(h-2i)$
 - base case, $n(1) = 1$; $n(2) = 2$; we take $i = h/2 - 1$ for $h \geq 3$
- ◆ Solving the base case we get: $n(h) > 2^{h/2-1}$ (see book p. 482/3)
- ◆ Taking logarithms: $h < 2 \log n(h) + 2$
- ◆ Thus the height of an AVL tree is $O(\log n)$

Insertion

- ◆ Insertion is as in a binary search tree
- ◆ Always done by expanding an external node.
- ◆ Example:




before insertion

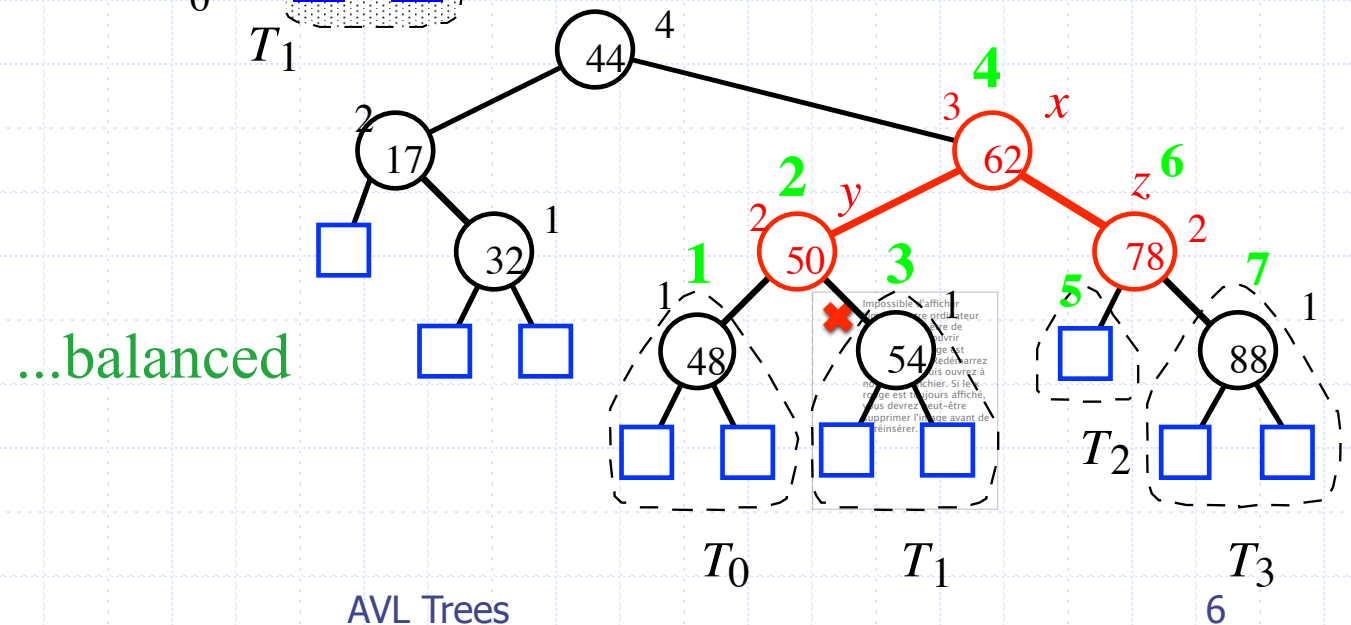
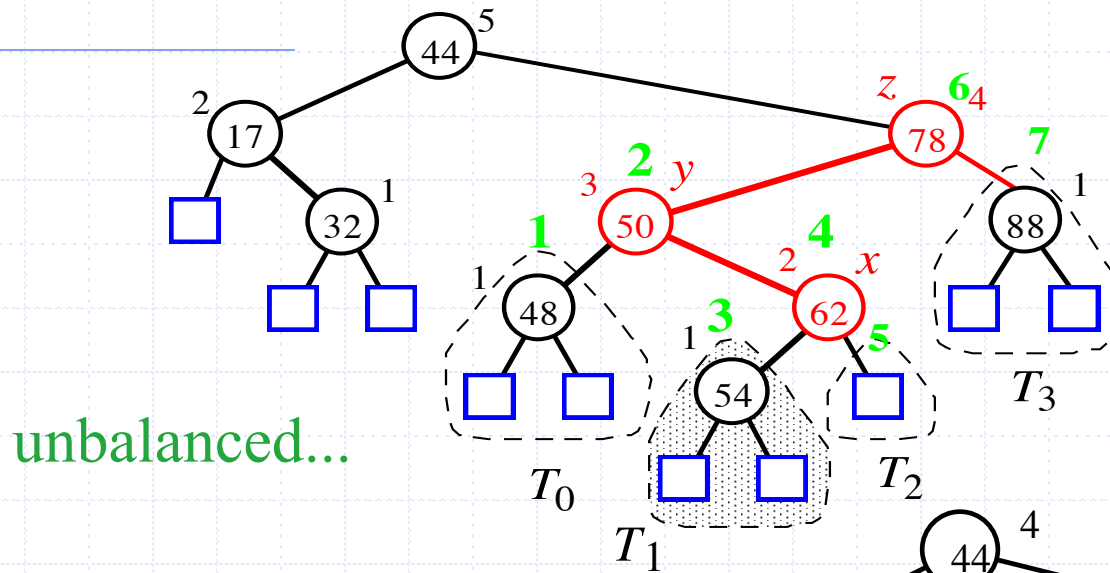
after insertion

Trinode Restructuring

- ◆ let (a, b, c) be an inorder listing of x, y, z
- ◆ perform the rotations needed to make b the topmost node of the three




Insertion Example, continued




Restructuring (as Single Rotations)

◆ Single Rotations:



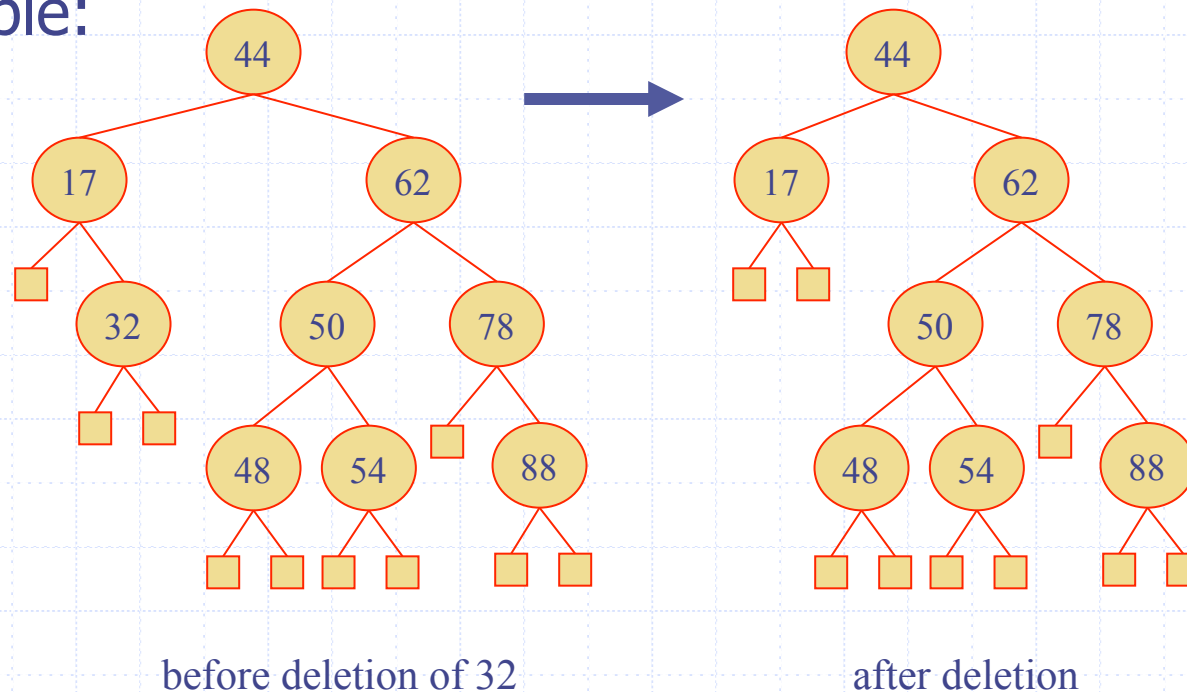
Restructuring (as Double Rotations)

◆ double rotations:




Removal

- ◆ Removal begins as in a binary search tree, which means the node removed will become an empty external node. Its parent, w, may cause an imbalance.
- ◆ Example:




Rebalancing after a Removal

- ◆ Let z be the **first unbalanced** node encountered while travelling up the tree from w . Also, let y be the child of z with the larger height, and let x be the child of y with the larger height
- ◆ We perform **restructure**(x) to restore balance at z
- ◆ As this restructuring may upset the balance of another node higher in the tree, we must continue checking for balance until the root of T is reached



AVL Tree Performance



- ◆ a single restructure takes $O(1)$ time
 - using a linked-structure binary tree
- ◆ Searching takes $O(\log n)$ time
 - height of tree is $O(\log n)$, no restructures needed
- ◆ Insertion takes $O(\log n)$ time
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$
- ◆ Removal takes $O(\log n)$ time
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$

Python Implementation

```
1 class AVLTreeMap(TreeMap):
2     """Sorted map implementation using an AVL tree."""
3
4     #----- nested _Node class -----
5     class _Node(TreeMap._Node):
6         """Node class for AVL maintains height value for balancing."""
7         __slots__ = '_height'          # additional data member to store height
8
9         def __init__(self, element, parent=None, left=None, right=None):
10             super().__init__(element, parent, left, right)
11             self._height = 0           # will be recomputed during balancing
12
13         def left_height(self):
14             return self._left._height if self._left is not None else 0
15
16         def right_height(self):
17             return self._right._height if self._right is not None else 0
```

Python Implementation, Part 2

```
18  #----- positional-based utility methods -----
19  def _recompute_height(self, p):
20      p._node._height = 1 + max(p._node.left_height(), p._node.right_height())
21
22  def _isbalanced(self, p):
23      return abs(p._node.left_height() - p._node.right_height()) <= 1
24
25  def _tall_child(self, p, favorleft=False): # parameter controls tiebreaker
26      if p._node.left_height() + (1 if favorleft else 0) > p._node.right_height():
27          return self.left(p)
28      else:
29          return self.right(p)
30
31  def _tall_grandchild(self, p):
32      child = self._tall_child(p)
33      # if child is on left, favor left grandchild; else favor right grandchild
34      alignment = (child == self.left(p))
35      return self._tall_child(child, alignment)
36
```

Python Implementation, end

```
37 def _rebalance(self, p):
38     while p is not None:
39         old_height = p._node._height           # trivially 0 if new node
40         if not self._isbalanced(p):             # imbalance detected!
41             # perform trinode restructuring, setting p to resulting root,
42             # and recompute new local heights after the restructuring
43             p = self._restructure(self._tall_grandchild(p))
44             self._recompute_height(self.left(p))
45             self._recompute_height(self.right(p))
46             self._recompute_height(p)           # adjust for recent changes
47         if p._node._height == old_height:       # has height changed?
48             p = None                             # no further changes needed
49         else:
50             p = self.parent(p)                   # repeat with parent
51
52     #----- override balancing hooks -----
53     def _rebalance_insert(self, p):
54         self._rebalance(p)
55
56     def _rebalance_delete(self, p):
57         self._rebalance(p)
```