



Red-Black Trees




From (2,4) to Red-Black Trees

- ◆ A red-black tree is a representation of a (2,4) tree by means of a binary tree whose nodes are colored **red** or **black**
- ◆ In comparison with its associated (2,4) tree, a red-black tree has
 - same logarithmic time performance
 - simpler implementation with a single node type



Red-Black Trees

- ◆ A red-black tree can also be defined as a binary search tree that satisfies the following properties:
 - **Root Property:** the root is black
 - **External Property:** every leaf is black
 - **Internal Property:** the children of a red node are black
 - **Depth Property:** all the leaves have the same black depth



class RedBlackTreeMap

```
class RedBlackTreeMap( TreeMap ):

class _Node( TreeMap._Node ):
    __slots__ = '_red'

    def __init__( self, element, parent = None, left = None, right = None ):
        super().__init__( element, parent, left, right )
        self._red = True

def _set_red( self, p ): p._node._red = True
def _set_black( self, p ): p._node._red = False
def _set_color( self, p, make_red ): p._node._red = make_red
def _is_red( self, p ): return p is not None and p._node._red
def _is_red_leaf( self, p ): return self._is_red( p ) and self.is_leaf( p )
```

Height of a Red-Black Tree

◆ **Theorem:** A red-black tree storing n items has height $O(\log n)$

Proof:

- The height of a red-black tree is at most twice the height of its associated (2,4) tree, which is $O(\log n)$
- ◆ The search algorithm for a binary search tree is the same as that for a binary search tree
- ◆ By the above theorem, searching in a red-black tree takes $O(\log n)$ time

Insertion (case root)

Root Property: the root is black

```
def _rebalance_insert( self, p ):
    #new node always red
    self._resolve_red( p )
```


```
def _resolve_red( self, p ):
    if self.is_root( p ):
        #make root black
        self._set_black( p )
    else:
```

...

Insertion

Internal Property: the children of a red node are black

- ◆ To insert (k, o) , we execute the insertion algorithm for binary search trees and color **red** the newly inserted node z unless it is the root
 - We preserve the root, external, and depth properties
 - If the parent v of z is black, we also preserve the internal property and we are done
 - Else (v is red) we have a **double red** (i.e., a violation of the internal property), which requires a reorganization of the tree
- ◆ Example where the insertion of 4 causes a double red:




Remedying a Double Red

- ◆ Consider a double red with child z and parent v , and let w be the sibling of v


Case 1: w is black

- The double red is an incorrect replacement of a 4-node
- **Restructuring**: we change the 4-node replacement




Case 2: w is red

- The double red corresponds to an overflow
- **Recoloring**: we perform the equivalent of a **split**




Restructuring

- ◆ A restructuring remedies a child-parent double red when the parent red node has a black sibling
- ◆ It is equivalent to restoring the correct replacement of a 4-node
- ◆ The internal property is restored and the other properties are preserved



Restructuring (cont.)


- ◆ There are four restructuring configurations depending on whether the double red nodes are left or right children




Insertion (case 1: black uncle)

else:

```
parent = self.parent( p )
if self._is_red( parent ):
    #double red problem
    uncle = self.sibling( parent )
    if not self._is_red( uncle ):
        #Case 1: misshapen 4-node
        middle = self._restructure( p ) #do trinode restructuring
        self._set_black( middle )      #and then fix colors
        self._set_red( self.left( middle ) )
        self._set_red( self.right( middle ) )
        #uncle was already black from the case
```




...



Recoloring

- ◆ A recoloring remedies a child-parent double red when the parent red node has a red sibling
- ◆ The parent v and its sibling w become black and the grandparent u becomes red, unless it is the root
- ◆ It is equivalent to performing a split on a 5-node
- ◆ The double red violation may propagate to the grandparent u



Insertion (case 2: red uncle)

else:

#Case 2: overflow (5-node)

grand = self.parent(parent)

self._set_red(grand)

self._set_black(self.left(grand))



self._set_black(self.right(grand))

self._resolve_red(grand)

#grandparent becomes red

#its children become black

#recur at red grandparent



Analysis of Insertion


Algorithm *insert*(k, o)

1. We search for key k to locate the insertion node z
2. We add the new entry (k, o) at node z and color z red
3. **while** *doubleRed*(z)
 if *isBlack*(*sibling*(*parent*(z)))
 $z = \text{restructure}(z)$
 return
 else { *sibling*(*parent*(z)) is red }
 $z = \text{recolor}(z)$

- ◆ Recall that a red-black tree has $O(\log n)$ height
- ◆ Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
- ◆ Step 2 takes $O(1)$ time
- ◆ Step 3 takes $O(\log n)$ time because we perform
 - $O(\log n)$ recolorings, each taking $O(1)$ time, and
 - at most one restructuring taking $O(1)$ time
- ◆ Thus, an insertion in a red-black tree takes $O(\log n)$ time

Deletion (the easy cases)

- ◆ To perform operation **remove(k)**, we first execute the deletion algorithm for binary search trees (results in the removal of a node that has at most one child, (either the node containing k or its inorder predecessor) and the promotion of its remaining child (if any)).
- ◆ Let v be the internal node removed, w the external node removed, and r the sibling of w
 - If either v or r was red, we color r black and we are done (corresponds to shrinking a 3-node or a 4-node).
 - When v black, then it either has zero children or it has one red leaf child (because the null subtree of the removed node has black height 0).
- ◆ Examples deletion of 3:




Deletion code

```
def _rebalance_delete( self, p ):
    #p is the parent of the node to be deleted (from TreeMap.py)
    if len( self ) == 1:
        #special case: ensure root is black
        self._set_black( self._root() )
    elif p is not None:
        n = self.num_children( p )
        if n == 1:
            #not a problem if red leaf, otherwise double black
            c = next( self.children( p ) )
            if not self._is_red_leaf( c ):
                self._fix_deficit( p, c )
        elif n == 2:
            #color black the promoted red leaf
            if self._is_red_leaf( self.left( p ) ):
                self._set_black( self.left( p ) )
            else:
                self._set_black( self.right( p ) )
```



Deletion (the double black case)

- ◆ Let v be the internal node removed, w the external node removed, and r the sibling of w
 - Else (v and r were both black) we color r **double black**, which is a violation of the internal property requiring a reorganization of the tree
WHY?
- ◆ Example where the deletion of 8 causes a double black:



Remedying a Double Black

- ◆ The algorithm for remedying a double black node w with sibling y




Remedying a Double Black

- ◆ The algorithm for remedying a double black node w with sibling y considers three cases

Case 1: y is black and has a red child

- We perform a **restructuring**, equivalent to a **transfer**, and we are done







```
def _fix_deficit( self, z, y ):
    #resolve black deficit at z, where y is the
    #root of z's heavier subtree
    if not self._is_red( y ):
        #y is black; will apply Case 1 or 2
        x = self._get_red_child( y )
        if x is not None:
            #Case 1: y is black and has
            #red child x; do "transfer"
            old_color = self._is_red( z )
            middle = self._restructure( x )
            self._set_color( middle, old_color )
            self._set_black( self.left( middle ) )
            self._set_black( self.right( middle ) )
        else:
            ...
```

Remedying a Double Black

- ◆ The algorithm for remedying a double black node w with sibling y considers three cases

Case 2: y is black and its children are both black

- We perform a **recoloring**, equivalent to a **fusion**, which may propagate up the double black violation



else:

```
#Case 2: y is black, but no red children;
#recolor as "fusion"
self._set_red( y )
if self._is_red( z ):
    self._set_black( z )
    #this resolves the problem
elif not self.is_root( z ):
    #recur upward
    self._fix_deficit( self.parent( z ),
                      self.sibling( z ) )
```

else:


...

Remedying a Double Black

- ◆ The algorithm for remedying a double black node w with sibling y considers three cases

Case 3: y is red

- We perform an **adjustment**, equivalent to choosing a different representation of a 3-node, after which either Case 1 or Case 2 applies



```
#Case 3: y is red; rotate misaligned
#3-node and repeat
self._rotate( y )
self._set_black( y )
self._set_red( z )
if z == self.right( y ):
    self._fix_deficit( z, self.left( z ) )
else:
    self._fix_deficit( z, self.right( z ) )
```

Remedying a Double Black

- ◆ The algorithm for remedying a double black node w with sibling y considers three cases
- ◆ Deletion in a red-black tree takes $O(\log n)$ time

Red-Black Tree Reorganization

Insertion

remedy double red

Red-black tree action

(2,4) tree action

result

restructuring

change of 4-node
representation

double red removed

recoloring

split

double red removed
or propagated up

Deletion

remedy double black

Red-black tree action

(2,4) tree action

result

restructuring

transfer

double black removed

recoloring

fusion

double black removed
or propagated up

adjustment

change of 3-node
representation

restructuring or
recoloring follows

Python Implementation

```
1 class RedBlackTreeMap(TreeMap):
2     """Sorted map implementation using a red-black tree."""
3     class _Node(TreeMap._Node):
4         """Node class for red-black tree maintains bit that denotes color."""
5         __slots__ = '_red'      # add additional data member to the Node class
6
7         def __init__(self, element, parent=None, left=None, right=None):
8             super().__init__(element, parent, left, right)
9             self._red = True    # new node red by default
```


Python Implementation, Part 2

```
10 #----- positional-based utility methods -----
11 # we consider a nonexistent child to be trivially black
12 def _set_red(self, p): p._node._red = True
13 def _set_black(self, p): p._node._red = False
14 def _set_color(self, p, make_red): p._node._red = make_red
15 def _is_red(self, p): return p is not None and p._node._red
16 def _is_red_leaf(self, p): return self._is_red(p) and self.is_leaf(p)
17
18 def _get_red_child(self, p):
19     """Return a red child of p (or None if no such child)."""
20     for child in (self.left(p), self.right(p)):
21         if self._is_red(child):
22             return child
23     return None
24
25 #----- support for insertions -----
26 def _rebalance_insert(self, p):
27     self._resolve_red(p) # new node is always red
28
29 def _resolve_red(self, p):
30     if self.is_root(p):
31         self._set_black(p) # make root black
32     else:
33         parent = self.parent(p)
34         if self._is_red(parent): # double red problem
35             uncle = self.sibling(parent)
36             if not self._is_red(uncle): # Case 1: misshapen 4-node
37                 middle = self._restructure(p) # do trinode restructuring
38                 self._set_black(middle) # and then fix colors
39                 self._set_red(self.left(middle))
40                 self._set_red(self.right(middle))
41             else: # Case 2: overfull 5-node
42                 grand = self.parent(parent)
43                 self._set_red(grand) # grandparent becomes red
44                 self._set_black(self.left(grand)) # its children become black
45                 self._set_black(self.right(grand))
46                 self._resolve_red(grand) # recur at red grandparent
```

Python Implementation, end

```
47 #----- support for deletions -----
48 def _rebalance_delete(self, p):
49     if len(self) == 1:
50         self._set_black(self.root()) # special case: ensure that root is black
51     elif p is not None:
52         n = self.num_children(p)
53         if n == 1: # deficit exists unless child is a red leaf
54             c = next(self.children(p))
55             if not self._is_red_leaf(c):
56                 self._fix_deficit(p, c)
57         elif n == 2: # removed black node with red child
58             if self._is_red_leaf(self.left(p)):
59                 self._set_black(self.left(p))
60             else:
61                 self._set_black(self.right(p))
62
63 def _fix_deficit(self, z, y):
64     """Resolve black deficit at z, where y is the root of z's heavier subtree."""
65     if not self._is_red(y): # y is black; will apply Case 1 or 2
66         x = self._get_red_child(y)
67         if x is not None: # Case 1: y is black and has red child x; do "transfer"
68             old_color = self._is_red(z)
69             middle = self._restructure(x)
70             self._set_color(middle, old_color) # middle gets old color of z
71             self._set_black(self.left(middle)) # children become black
72             self._set_black(self.right(middle))
73         else: # Case 2: y is black, but no red children; recolor as "fusion"
74             self._set_red(y)
75             if self._is_red(z):
76                 self._set_black(z) # this resolves the problem
77             elif not self.is_root(z):
78                 self._fix_deficit(self.parent(z), self.sibling(z)) # recur upward
79         else: # Case 3: y is red; rotate misaligned 3-node and repeat
80             self._rotate(y)
81             self._set_black(y)
82             self._set_red(z)
83             if z == self.right(y):
84                 self._fix_deficit(z, self.left(z))
85             else:
86                 self._fix_deficit(z, self.right(z))
```